

How Augmentible Began: From Overwhelm to Augmentation

Ten months ago, I opened ChatGPT for something very specific and very small: I wanted help understanding why I kept getting physics questions wrong for my FRCR Part 1 Physics. Not to cheat, not to shortcut—just to *see* where my thinking was off. I pasted in a few practice questions, asked it to walk me through the wrong answers step-by-step, and within minutes I had clearer explanations than I'd given myself in hours of revision. That small win planted a seed: maybe this wasn't a gimmick; maybe it was a tutor with infinite patience.

Not too long after passing the exam, life was getting full. I was preparing for on-calls, tracking endless to-dos, dealing with a summer calendar full of friends' weddings, and staring down a wall of textbooks that all felt important and somehow equally unapproachable. I wasn't short on motivation—I was short on structure and time. The work was doable; the *orchestration* wasn't.

A friend then mentioned "ChatGPT Projects." He'd uploaded PDF copies of his key notes and textbooks, and was using targeted prompts: "Explain this CT finding with reference to *[textbook]*, give me the page, then cross-check with Radiopaedia." It sounded a bit like magic. I tried it. It wasn't magic—it was better. It was *repeatable*.

Suddenly, my study sessions shifted. Instead of flipping between chapters, bookmarks, and browser tabs, I could ask for a focused explanation and get cited, cross-referenced guidance I could verify. I wasn't replacing the textbooks; I was *triaging* them. The time saved didn't just feel efficient; it felt humane. Space returned—for proper thinking, for rest, for the other parts of life that do not fit neatly into a rota.

That led to a quiet, ongoing experiment: if AI can help with this, what else can it help with—without lowering the bar on accuracy or clinical judgement? I started building small, then iterating:

- Summarising journal articles to decide whether they deserved a full read.
- Turning messy bullet points into teaching slides, faster.
- Logging learning points and action items from meetings in a way I'd actually revisit.

Each improvement on its own was modest. Together, they changed the *feel* of work. Less scrambling, more clarity. Less context switching, more continuity.

Of course, there was a wobble. Like many clinicians first encountering this kind of technology, I worried about job security and professional identity. If AI can draft, summarise, and pattern-match, where do we fit? That anxiety is understandable—and healthy. But it misses the practical truth I've seen again and again: AI +

professional beats AI alone or professional alone. The model is fast; the clinician is wise. The model is broad; the clinician is accountable. The model spots a pattern; the clinician knows which pattern matters *here, now, for this person*. Augmentation isn't a slogan—it's the only arrangement that makes sense for safe, humane care.

This is the mindset behind **Augmentible**. It wasn't born from techno-optimism; it came from feeling stretched and deciding to experiment my way out of it. We're building a place where healthcare professionals can learn the *applied* stuff: the prompts, the workflows, the governance guardrails, and the very human habits that let you use these tools without losing your judgement (or your weekend!)

Right now, we're exploring a set of everyday helpers that smooth out the friction in clinical and academic life:

- **AI meeting notes** that capture decisions, owners, and deadlines you'll actually follow.
- **AI-assisted calendar planning** that translates goals into realistic time blocks (and respects clinics and on-calls).
- **Lightweight automations** for admin you'd prefer not to spend brainpower on—think templated emails, reading queues, audit logs, and tidy archives that don't become digital attics.

We'll share the playbooks openly—what works, what doesn't, and why. Expect step-by-step guides, short videos, and templates you can copy, plus honest notes about failure modes (because every "automation" breaks somewhere, and it helps to know where).

If you're new to this, here's the simplest way to begin:

1. **Pick one friction point** (e.g., preparing for an exam, summarising a paper, or structuring teaching).
2. **Write down your current steps**, then ask an AI tool to draft the same output from your inputs.
3. **Keep your judgement in the loop.** Compare, edit, and decide. Save the working prompt as a mini-template.
4. **Repeat next week** for a second task. Tiny, boring wins compound.

Augmentible is my attempt to bottle those tiny wins and hand them to colleagues who are as busy—and as ambitious—as I was when this started. Not to replace hard work, but to aim it better.

If this resonates, I'd love you to be part of the community. We're putting together a newsletter with practical how-tos, prompts that survive the real world, and occasional reflections like this one. You can reach me directly at

yogeshwar@augmentible.com—questions, use-cases, pushback, all welcome. And if you want the next post when it drops, **subscribe to the newsletter** so we can send it straight to you.

Augmentation over exhaustion. That's the project. Let's see how far it goes—
together.

Yogi